



Data
Schools

Module 1:

Introduction to Machine Learning

Overview

Faculty Introduction

Ekpe Okorafor, PhD



AFFILIATIONS:

- Accenture - Applied Artificial Intelligence
- African University of Science & Technology
- Nigerian British University

RESEARCH INTERESTS:

- Generative AI
- Machine Learning
- Big Data & Advanced Analytics
- Knowledge Graph
- Information Assurance and Cybersecurity

FUN & INTERESTS

- Drawing & Painting
- Hiking
- Soccer
- DJ & Dancing

Modules for this topic

1. Overview: What is Machine learning
2. Categories of Machine Learning
3. Machine Learning Application Development Approach
4. *Building Classification Model*
5. Recommender Systems
6. *Building a Recommender Engine*

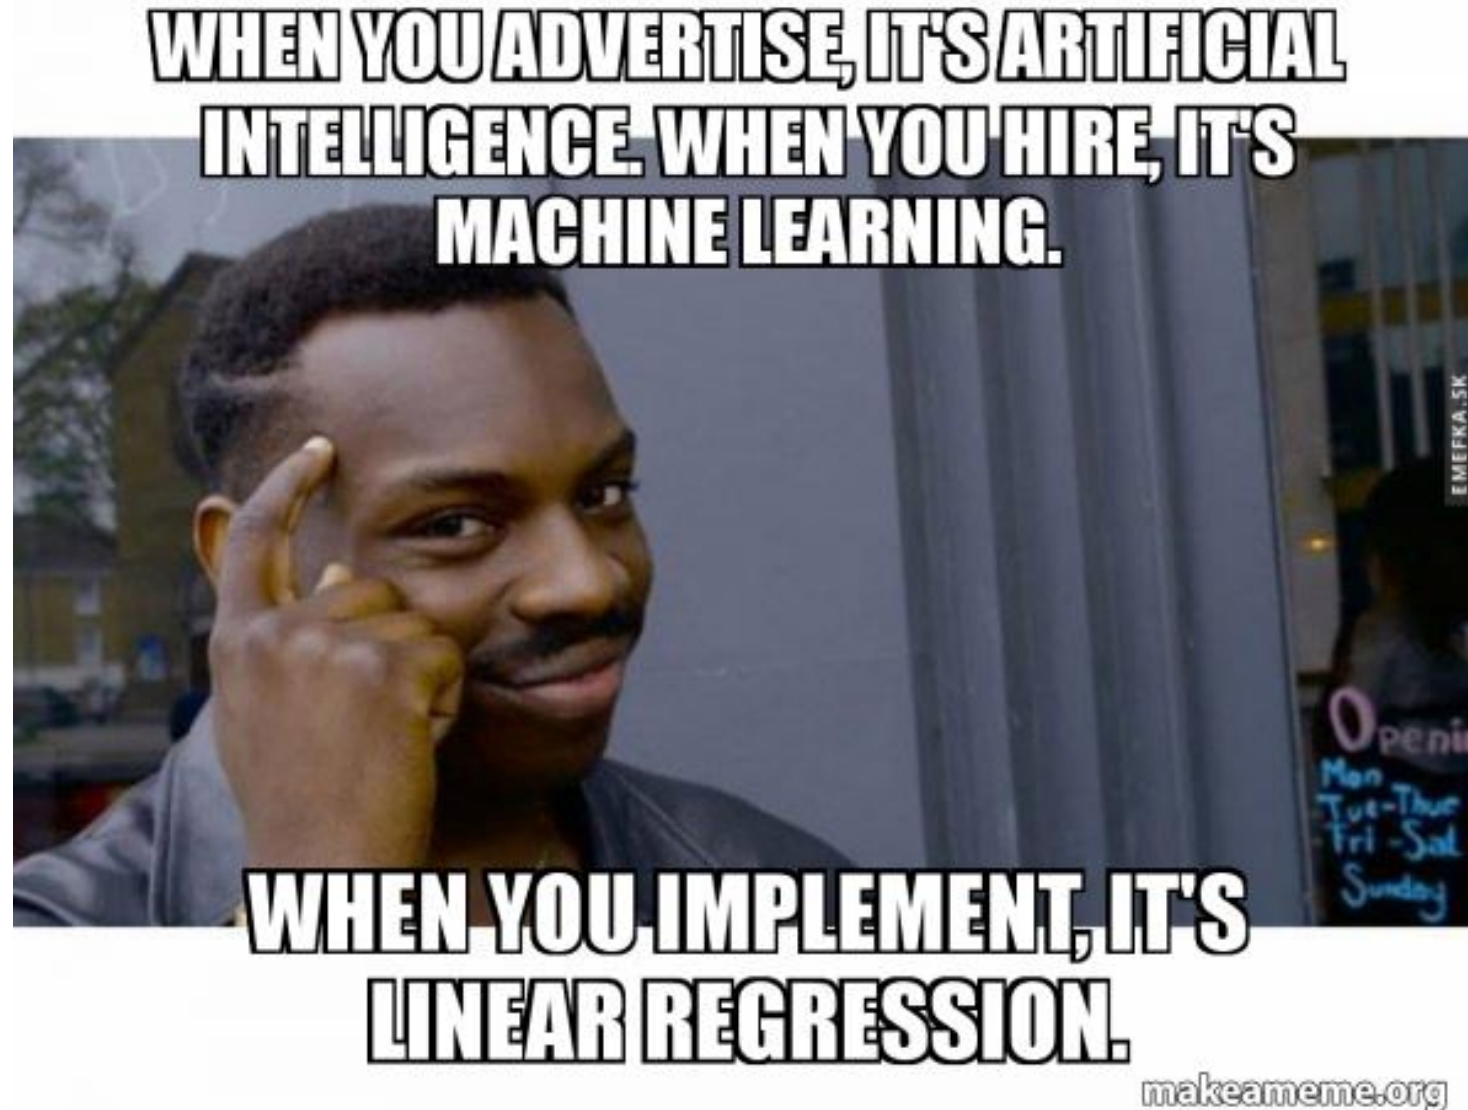
Learning Objectives

Upon successful completion of this topic, you will be able to:

- Define machine learning
- Describe the categories of machine learning
- Decide when to leverage Machine learning
- Build a simple classifier model
- Discuss approaches to ML application development
- Differentiate between the ML approaches and motivations
- Build a simple recommender engine in R

Module 1

Overview: What is Machine Learning?



Why Do We Want Machines To Learn?

- Lucia wants to buy a car
 - New car = \$20,000
 - 1 yr old (used) = \$19,000
 - 2 yr old (used) = \$18,000 and so on
- Pattern -> car price depends on its age and drops \$1,000 every year but won't get lower than \$10,000.
- Lucia invented *regression*
- **Problem:** There are so many other factors, e.g., options, technical condition, seasonal demand spikes, etc.
- **Challenge:** Lucia can't keep all that data in her head while calculating the price.
- **Solution:** Provide the machine some data and ask it to find all hidden patterns related to price.


What is What is Machine Learning?

```
100100011101000000101000110111010110
10010011110111000000111100110100100
100001101101111101010011100001101001
111111010000110111001010111100001011
1100111110111111100100001110110110
010000110100110110000110000100010000
010101110011001111011001110100010111
001000010101100101000001000010011110
011101001111100101110101010111100
100010000101100010101101010111000101
010010000100101011110011100001010000
010110000010011101010010101110110001
01101111101011100010100010100010000
011010011011011010001000101111001101
000101000001100110001100100010010110
10010101010001001110010101010111101
```

Data

$$f_x$$


Features



Algorithms

“Machine learning is the field of study that gives computers the ability to learn without being explicitly programmed. – Arthur L. Samuel, AI Pioneer, 1959

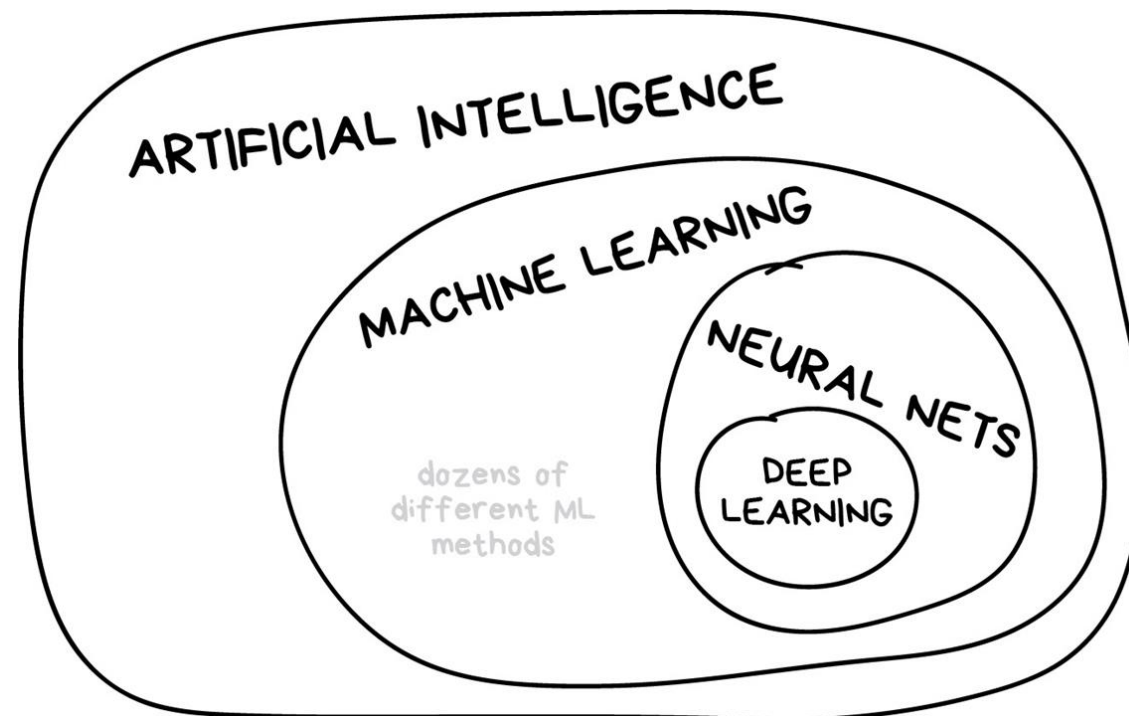
What is Machine Learning?



“Machine learning is the field of study that gives computers the ability to learn without being explicitly programmed.” – Arthur L. Samuel, AI Pioneer, 1959

What is Machine Learning?

- Artificial Intelligence (AI) is a branch of Computer Science that uses algorithms and techniques to mimic human intelligence
- Machine Learning (ML) is one of several AI techniques for sophisticated cognitive tasks



Machine Learning

Machine Learning is a particularly interesting technique because it represents a paradigm shift within AI

Traditional AI techniques



- **Static** – hard-coded set of steps and scenarios
- **Rule Based** – expert knowledge
- **No generalization** – handling special cases is difficult

Machine Learning



- **Dynamic** – evolves with data, finds new patterns
- **Data driven** – discovers knowledge
- **Generalization** – adapts to new situations and special cases

Machine Learning - Example

Example - Excelling at playing the game of chess



Symbolic AI

“Let us sit down with the world’s best chess player, Ekpe Okorafor, and put his knowledge into a computer program”

Mathematical/Statistical AI

“Let us simulate all the different possible moves and the associated outcomes at each single step and go with the most likely to win”

Machine Learning Approach

“Let us show millions of examples or real life and simulated games (won and lost) to the program, and let it learn from experience”

Machine Learning – When to use

Machine learning is particularly good at solving 2 types of problems where other AI techniques fail

Tasks programmers can't describe

Handwriting



Cognitive Reasoning

Complex multidimensional problems that can't be solved by numerical reasoning

Weather Forecasting



Health Care Outcomes



Network Intrusion



Movie Recommendation

Applications of Machine Learning


- Email spam detection
- Face detection and matching (e.g., iPhone X, Windows laptops, etc.)
- Web search (e.g., DuckDuckGo, Bing, Baidu, Google)
- Sports predictions
- Post office (e.g., sorting letters by zip codes)
- ATMs (e.g., reading checks)
- Credit card fraud
- Stock predictions
- Smart assistants (Apple Siri, Amazon Alexa, . . .)
- Product recommendations (e.g., Walmart, Netflix, Amazon)
- Self-driving cars (e.g., Uber, Tesla)
- Language translation (Google translate)
- Sentiment analysis
- Drug design
- Medical diagnoses
-

Exercise

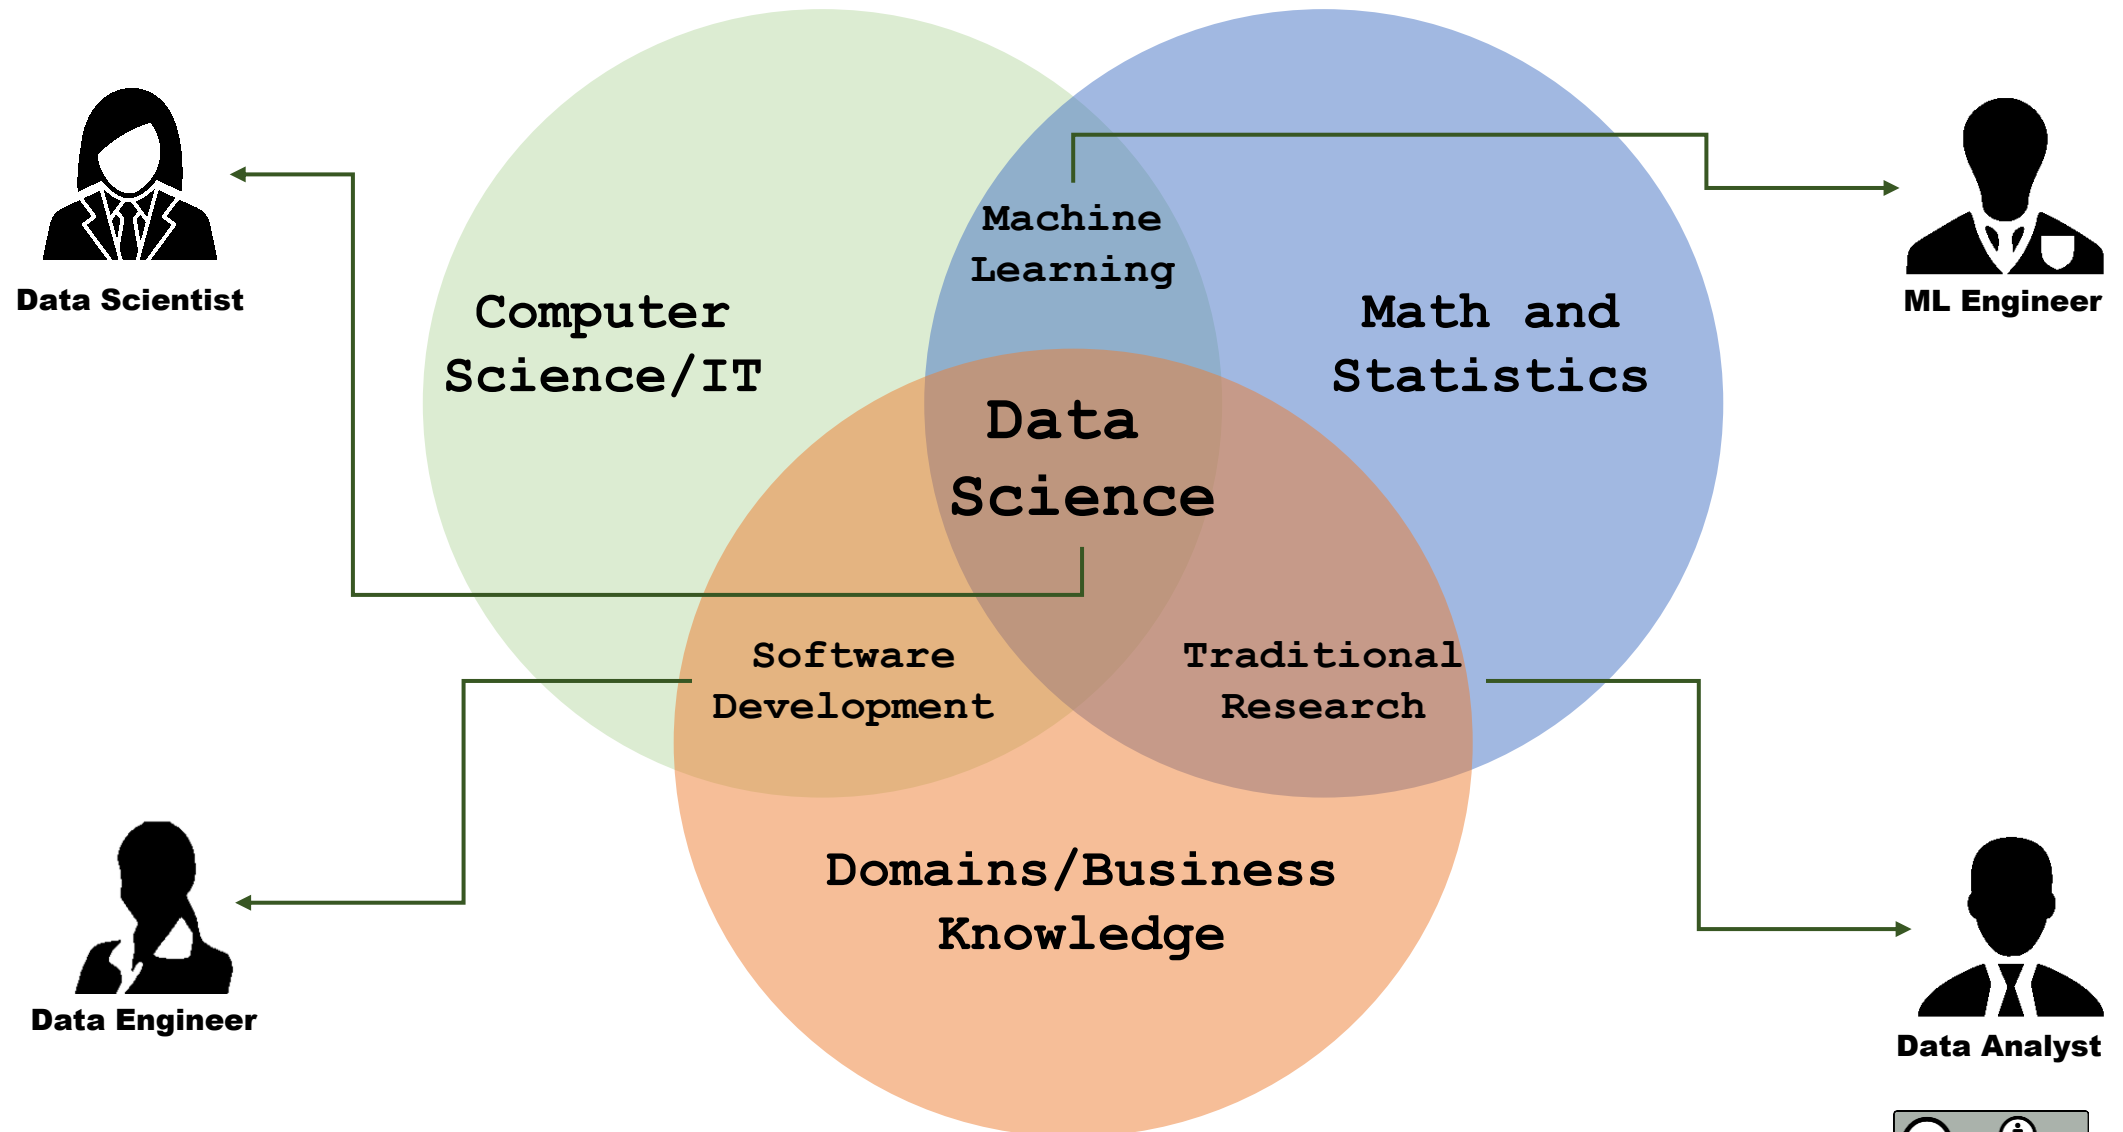
As a first exercise, think about how machine learning could be applied in these problem areas or tasks listed in the previous slide:

- What is the desired outcome?
- What could the dataset look like?
- Is this a supervised or unsupervised problem, and what algorithms would you use? (Supervised and unsupervised learning will be introduced next)
- How would you measure success?
- What are potential challenges or pitfalls?

Class Discussion



The World of Data



Summary

- We now have an overview and a good definition for Machine Learning
- We have seen some examples where ML can be leveraged