

---

# Cyber Infrastructures

## Cloud Computing

Instructor: Raphael Cobe  
@[raphaelmcobe](https://twitter.com/raphaelmcobe) / [raphaelmcobe@gmail.com](mailto:raphaelmcobe@gmail.com)

# About the Instructor



# About the Instructor

- ▷ Researcher at the [Advanced Institute for Artificial Intelligence](#);
- ▷ Associate Researcher at the [São Paulo State University - UNESP](#);
- ▷ Member of the [SPRACE project](#) (a [CMS LHC collaboration](#));
- ▷ Chair of the CODATA-RDA Data School groups since 2018;
- ▷ Background on Artificial Intelligence and High Performance Computing;

# Materials

## Exercises for the Lecture

1. [Warmup/Data Exploration using Google Colab;](#)
2. [Regression Analysis using Apache Spark on Colab;](#)
3. [Setting up a cloud resource using OpenStack;](#)
4. [Accessing a Virtual Machine using the SSH protocol;](#)
5. [Deploying a Jupyter Notebook Server at an IaaS platform using the conda environment manager;](#)
6. [Deploying a Jupyter Notebook Server at an IaaS platform using Docker containers](#)

# Acknowledgements 🙌

The people that worked hard to get this materials/resources available

1. Rob Quick (Indiana University) for working on the exercises;
2. [CODATA-RDA Schools of Research Data Science](#)  
Co-chairs for reviewing and evaluating the materials;

# The setup problem

- ▷ Your science computing is **too complex!**
  - Monte carlo simulation, image analysis, genetic algorithm, ...
- ▷ It will take **a year (CPU time)** to get the results on your laptop, but your **paper is due in a week.**
- ▷ What do you do?

**Option 1:** Wait a year

**Option 2:** Local Clustered Computing

**Option 3:** Use a High Performance  
Computing(HPC)

**Option 4:** Use lots of commodity computers

**Option 5:** Buy (or Borrow) some computing  
from a Cloud Provider

# These are All Valid Options

- ▷ Remember the problem you have one month to publish results for your conference
  - **Option 1:** You will **miss your deadline**
  - **Option 2:** You **might miss your deadline** – But if you're lucky you'll make it (or if you know the admin)
  - **Option 3:** If you have parallelized code and can get an allocation you **have a good chance**
  - **Option 4:** If you can serialize your workflow you **have a good chance**
  - **Option 5:** **You can meet your deadline** for a price. Though **academic clouds are becoming more available.**



**Option 1:** Wait a year

**Option 2:** Local Clustered Computing

**Option 3:** Use a High Performance  
Computing(HPC)

**Option 4:** Use lots of commodity computers

**Option 5:** Buy (or Borrow) some computing  
from a Cloud Provider

# Computing Infrastructures

- ▷ **Local Laptop/Desktop** – Short jobs with small data
- ▷ **Local Cluster** – Larger jobs and larger data but subject to availability
- ▷ **HPC** – Prime performance with parallelized and optimized code
- ▷ **HTC** – Sustained computing over a long period for serialized
- ▷ **Cloud** – Need deeper permission on an OS and have deeper pockets

# Agenda

- ▷ Introduction/motivation;
- ▷ A few words about virtualization;
- ▷ Cloud environments;
  - Public vs Private vs Hybrid
  - Cloud deployment models: SaaS, PaaS, and IaaS;
- ▷ Jupyter Notebooks;
- ▷ Openstack cloud manager;
- ▷ Other virtualization techniques;

# Agenda

- ▷ Introduction/motivation;
- ▷ A few words about virtualization;
- ▷ Cloud environments;
  - Public vs Private vs Hybrid
  - Cloud deployment models: SaaS, PaaS, and IaaS;
- ▷ Jupyter Notebooks;
- ▷ Openstack cloud manager;
- ▷ Other virtualization techniques;

# Introduction

## Cloud Computing

- ▷ Complete solution in which all the features of computing are quickly provided to users as the demand increases;
- ▷ You can control the features and services offered to ensure high availability, security and quality of service
  - Users get the resources they need, no more, no less.

# Introduction

## Cloud Computing - Some history

- ▷ **Utility Computing** - John McCarthy
- ▷ ARPANET (Advanced Research Projects Agency Network) - 1969
  - “Intergalactic Computer Network” in which everyone on the planet would be **interconnected by way of computers**;
- ▷ Cloud computing concepts: **availability** and **accessibility**

# Introduction



# Introduction

## Cloud Computing

- ▷ Creates an idea of availability of **infinite resources**;
- ▷ Eliminates the need to acquire and provision resources in advance;
- ▷ Offers **elasticity**, allowing companies to use resources as needed;
- ▷ Payment for services is made by the **amount of resources used**;
- ▷ Deeply connected to the **Virtualization** concept;



# Agenda

- ▷ Introduction/motivation;
- ▷ **A few words about virtualization;**
- ▷ Cloud environments;
  - Public vs Private vs Hybrid
  - Cloud deployment models: SaaS, PaaS, and IaaS;
- ▷ Jupyter Notebooks;
- ▷ Openstack cloud manager;
- ▷ Other virtualization techniques;

# Introduction

## Virtualization

- ▷ Simulated or virtual computing environment instead of a physical environment;
- ▷ Include computer-generated versions of hardware, storage devices, and others;
- ▷ This allows organizations to partition a single physical computer or server into multiple virtual machines;
  - Interact independently and run different operating systems;
  - Share the resources of a single host computer;

# Introduction

## Virtual Machines

- ▷ It has CPU, memory, disks to store your files and can connect to the Internet if needed;
- ▷ Software-defined computers on physical servers, existing only as code;
  - "Virtual" version of a computer, with dedicated amounts of CPU, memory, and storage that are "borrowed" from a physical host computer;
- ▷ The virtual machine is partitioned from the rest of the system;

# Introduction

## Virtual Machines - Motivation

- ▷ Create and **deploy applications in the cloud**;
- ▷ **Try a new OS** (operating system), including beta versions;
- ▷ **Create a new environment** to make running development and test scenarios simpler and faster for developers;
- ▷ **Back up** your existing operating system;

# Introduction

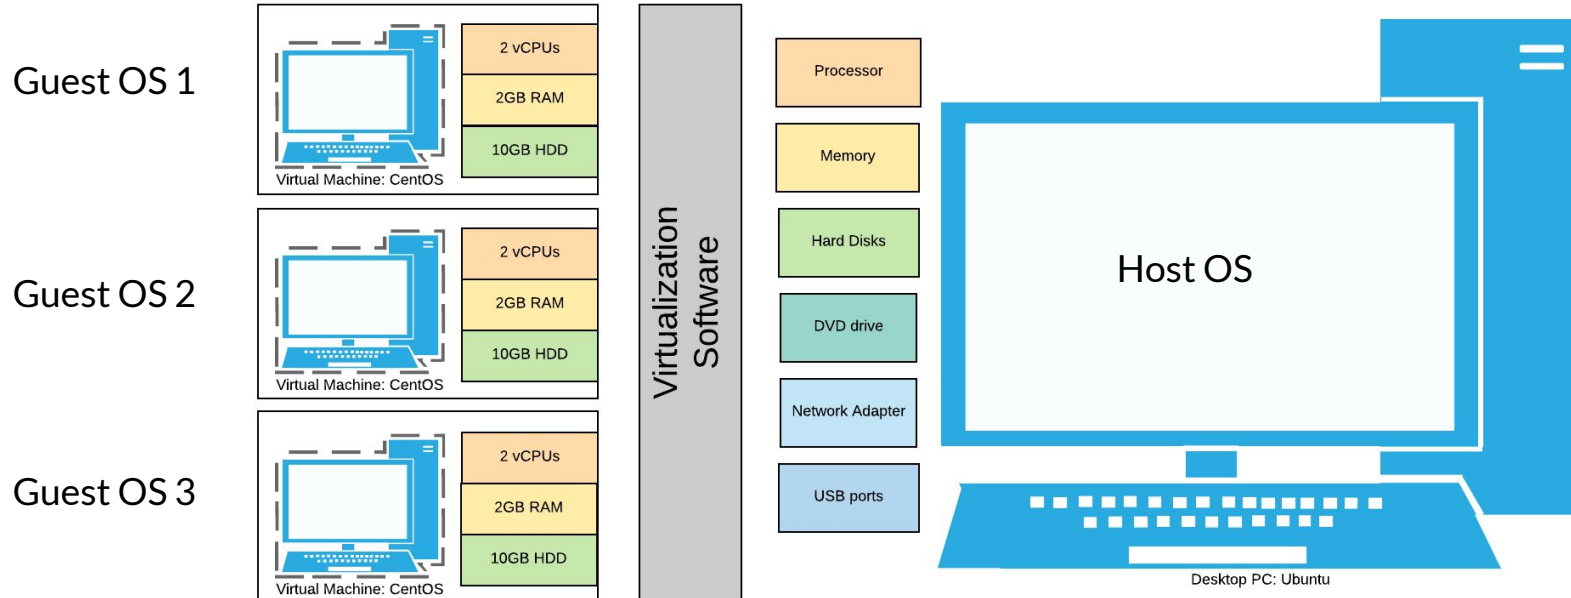
## Virtual Machines - Benefits

- ▷ Remain **completely independent** of each other and of **the physical host computer**;
- ▷ **Hypervisor** - allows you to run different operating systems on different virtual machines at the same time;
  - E.g., running a windows virtual machine inside a linux laptop;
- ▷ You can **move a VM from one hypervisor to another completely different computer** almost instantly;

# Introduction

## Virtual Machines

### Hardware Virtualization: a Desktop Virtualization Example



# Introduction

## Hypervisors - How does it work?

- ▷ Can be **run on an operating system** (such as a laptop) or **installed directly on hardware** (such as a server);
- ▷ **Translates operations** on virtual HW into host OS instructions;
  - The host OS then runs the instructions on real HW;
- ▷ Can cause **performance loss**;

# Agenda

- ▷ Introduction/motivation;
- ▷ A few words about virtualization;
- ▷ **Cloud environments;**
  - Public vs Private vs Hybrid
  - Cloud deployment models: SaaS, PaaS, and IaaS;
- ▷ Jupyter Notebooks;
- ▷ Openstack cloud manager;
- ▷ Other virtualization techniques;



# Cloud Computing

## Cloud Environments

- ▷ Abstract, group and share scalable resources across a network;
- ▷ Not a technology per se;
- ▷ Provide a self-service interface;
- ▷ Elasticity and the ability to scale up and down;
- ▷ Application programming interfaces (APIs);
- ▷ Billing and metering of service usage in a pay-as-you-go model;

# Cloud Computing

## Cloud Environments - Public, Private and Hybrid

- ▷ **Public:** A cloud environment **built from resources without an owner**, such as an end user, that can be redistributed to other tenants.
- ▷ **Private:** Deployed using **local resources**. It is sometimes preferred for its **ability to offer dedicated resources**;
- ▷ **Hybrid:** **Connects web resources and current resources not in the cloud**. Allows to extend and grow a company's cloud infrastructure;

# Cloud Computing

## Cloud Deployment Models - Infrastructure As A Service

- ▷ Gives you the **highest level of flexibility** and management control over your IT resources;
- ▷ Users **control the applications, data, operating system, and execution environments**;
- ▷ The IaaS provider provides the **virtualization, storage, networking and servers**.
- ▷ The user **does not need to have an on-premise datacenter** or worry about maintenance or physical updates;

# Cloud Computing

## Cloud Deployment Models - Platform As A Service

- ▷ Application hardware and software platforms are **provided by third parties**;
- ▷ Mainly **aimed at developers and programmers**;
  - Allows the user to **develop, run and manage applications**;
  - A **provider hosts the hardware and software components** in its infrastructure
- ▷ Provides **a platform** to the user as an integrated solution;

# Cloud Computing

## Cloud Deployment Models - Software As A Service

- ▷ Provides users with a **cloud application with the IT infrastructure and platforms** underlying it;
- ▷ Compatible with **software subscription models**;
- ▷ Service performance is **determined by internet connection speed**;
- ▷ **General public** oriented:
  - Does not require IT/Software development skills;
- ▷ Allows providers to **easily deploy new features** to customers.

# Cloud Computing

## Cloud Deployment Models

- Providers:



# Agenda

- ▷ Introduction/motivation;
- ▷ A few words about virtualization;
- ▷ Cloud environments;
  - Public vs Private vs Hybrid
  - Cloud deployment models: SaaS, PaaS, and IaaS;
- ▷ **Jupyter Notebooks;**
- ▷ Openstack cloud manager;
- ▷ Other virtualization techniques;

# PaaS Example

## Jupyter Notebooks

- ▷ Contain **both computer code and rich text elements** (equations, figures, links, etc.);
- ▷ Both **human-readable and executable documents**;
  - Contain the **analysis description and the results** (figures, tables, etc..) as well as **code that can be run** to perform data analysis;
- ▷ "*Jupyter*" is an **acronym for Julia, Python and R**;
- ▷ **Replication and reproducibility** are two pillars of the scientific method.
  - Jupyter notebooks make it easy to do both;



# PaaS Example

## Jupyter Notebooks history



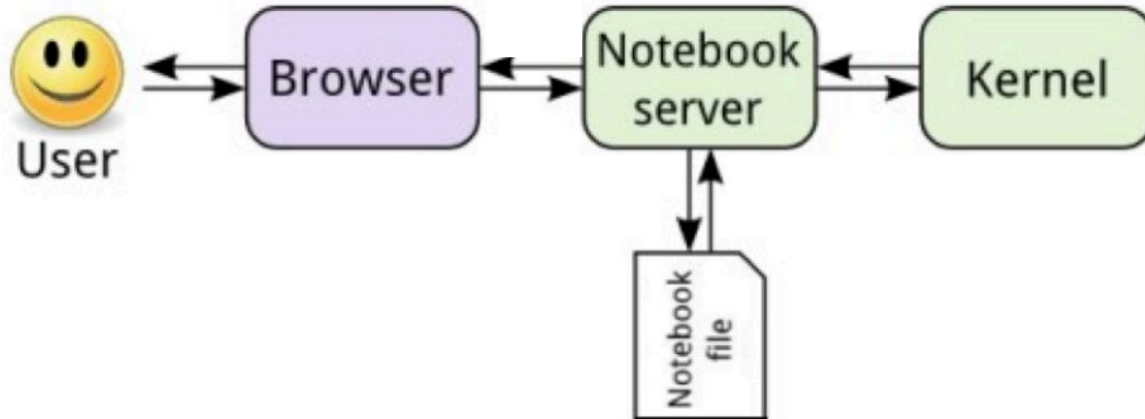
# PaaS Example

## Jupyter Notebooks - concepts

- ▷ **Cells:**
  - **Blocks that can be executed** individually and present results;
  - **Code cell:** contains **code to be executed in the kernel** and displays its output below;
  - **Markdown cell:** contains **text formatted using Markdown** and displays its output in place
- ▷ **Kernels:**
  - Computational engine that runs the code contained in a notebook.

# PaaS Example

Jupyter Notebooks - how does it work?



# PaaS Example

## Jupyter Notebooks - Google Colaboratory



- ▷ Free cloud service hosted by Google
- ▷ Encourage research on Machine Learning and Data Science;
- ▷ Jupyter notebook environment that requires no setup and runs in the cloud;

# PaaS Example

## Jupyter Notebooks - Google Colaboratory

- ▷ Compatible with codes in **Python** (2.7 and 3.6) and **R** 4.1;
- ▷ Access to **powerful scientific computing resources**;
- ▷ **Free GPU** acceleration;
- ▷ **Pre-installed** all major libraries
- ▷ **Collaboration** works like Google Docs;
- ▷ **Supports bash** commands;
- ▷ **Stored in Google Drive**.

# PaaS Example

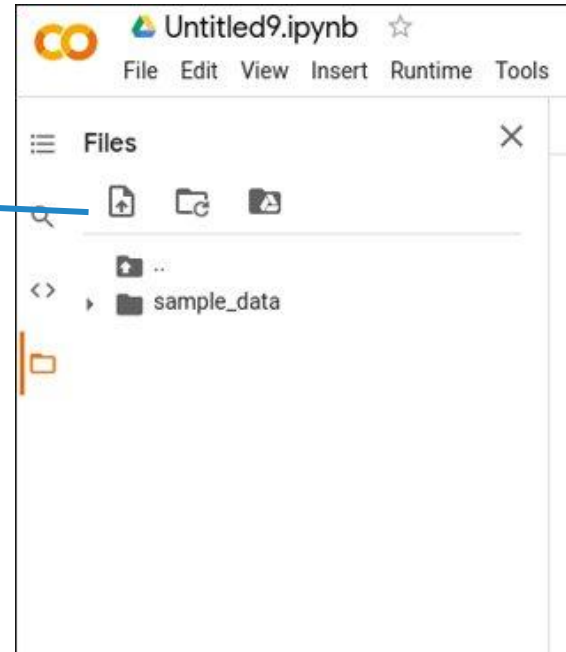
## Jupyter Notebooks - Google Colaboratory

- ▷ How to access:
  - <https://colab.to> (Python) or
  - <https://colab.to/r> (R)
- ▷ Bash commands:
  - Prefixing the command with “!” (Python)
  - Use **`system("COMMAND", intern=T)`** (R)
- ▷ Installing new libraries:
  - pip support (!pip install [package name])
  - In R: `install.packages('[package name]')`

# PaaS Example

## Jupyter Notebooks - Google Colaboratory

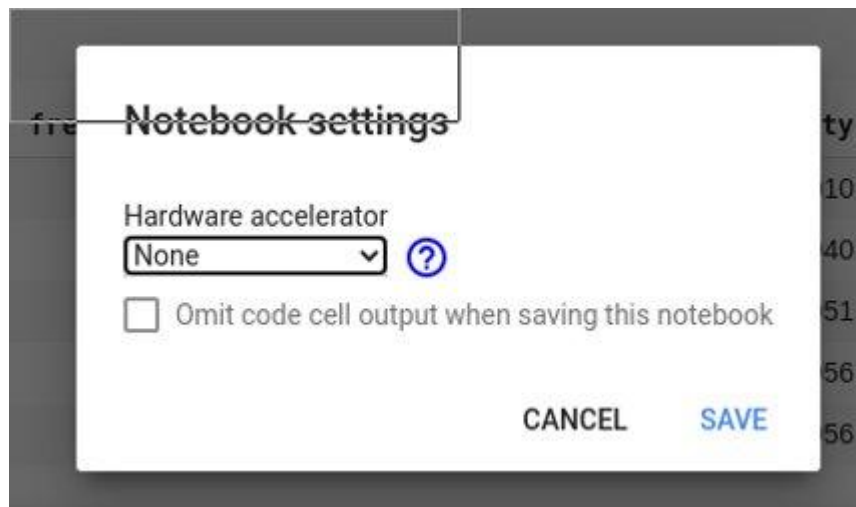
- ▷ Uploading data:



# PaaS Example

## Jupyter Notebooks - Google Colaboratory

- ▶ Choosing the Runtime:
  - CPU, GPU ou TPU (limited to 12 hours/day);





# Agenda

- ▷ Introduction/motivation;
- ▷ A few words about virtualization;
- ▷ Cloud environments;
  - Public vs Private vs Hybrid
  - Cloud deployment models: SaaS, PaaS, and IaaS;
- ▷ Jupyter Notebooks;
- ▷ **Openstack cloud manager;**
- ▷ Other virtualization techniques;

# Openstack

## Some history

2010

NASA + Rackspace  
develop the basis of  
OpenStack

2014

OpenStack Marketplace opens to  
showcase maturing ecosystem;  
“Juno” release seen as enterprise grade

2016 - April

Half the Fortune 100 run  
OpenStack; Certified OpenStack  
Administrator program launched

2017

OpenStack emerges as one  
platform for containers,  
VMs and bare metal

2012

OpenStack Foundation  
established

2015

OpenStack Powered  
interop certification  
launched

2016

China booms; 86% of  
telecoms say OpenStack  
important to their business

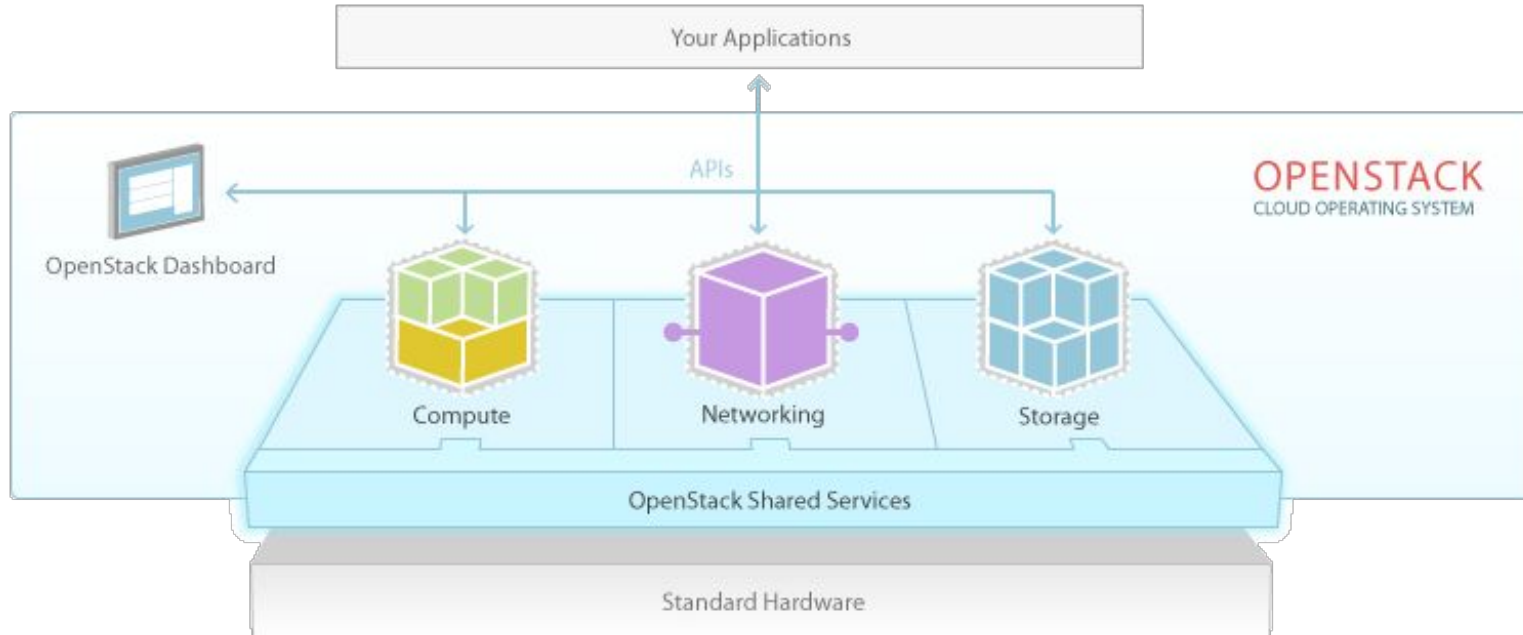
# Openstack

## Introduction

- ▷ Opensource;
- ▷ Large project with several components;
  - Virtual Machines (*nova*), Virtual Networks (*neutron*), Cloud Storage (*cinder, swift, glance*) User self-service interface (*horizon*), User authentication and authorization (*keystone*), etc.;
- ▷ OpenStack does not virtualize resources alone, but uses them to create clouds;

# Openstack

## What is in it?



# Openstack

## Main concepts

- ▷ **Project**: A container used to group a set of resources such as virtual machines, volumes and images with the same access rights and quota.
- ▷ **Quota**: A per-project limit such as the total number of cores or RAM permitted for a set of virtual machines.
- ▷ **Flavor**: The definition of the size of a virtual machine and its characteristics
  - E.g., 2 core virtual machine with 8 GB of RAM

# Openstack

## Main concepts

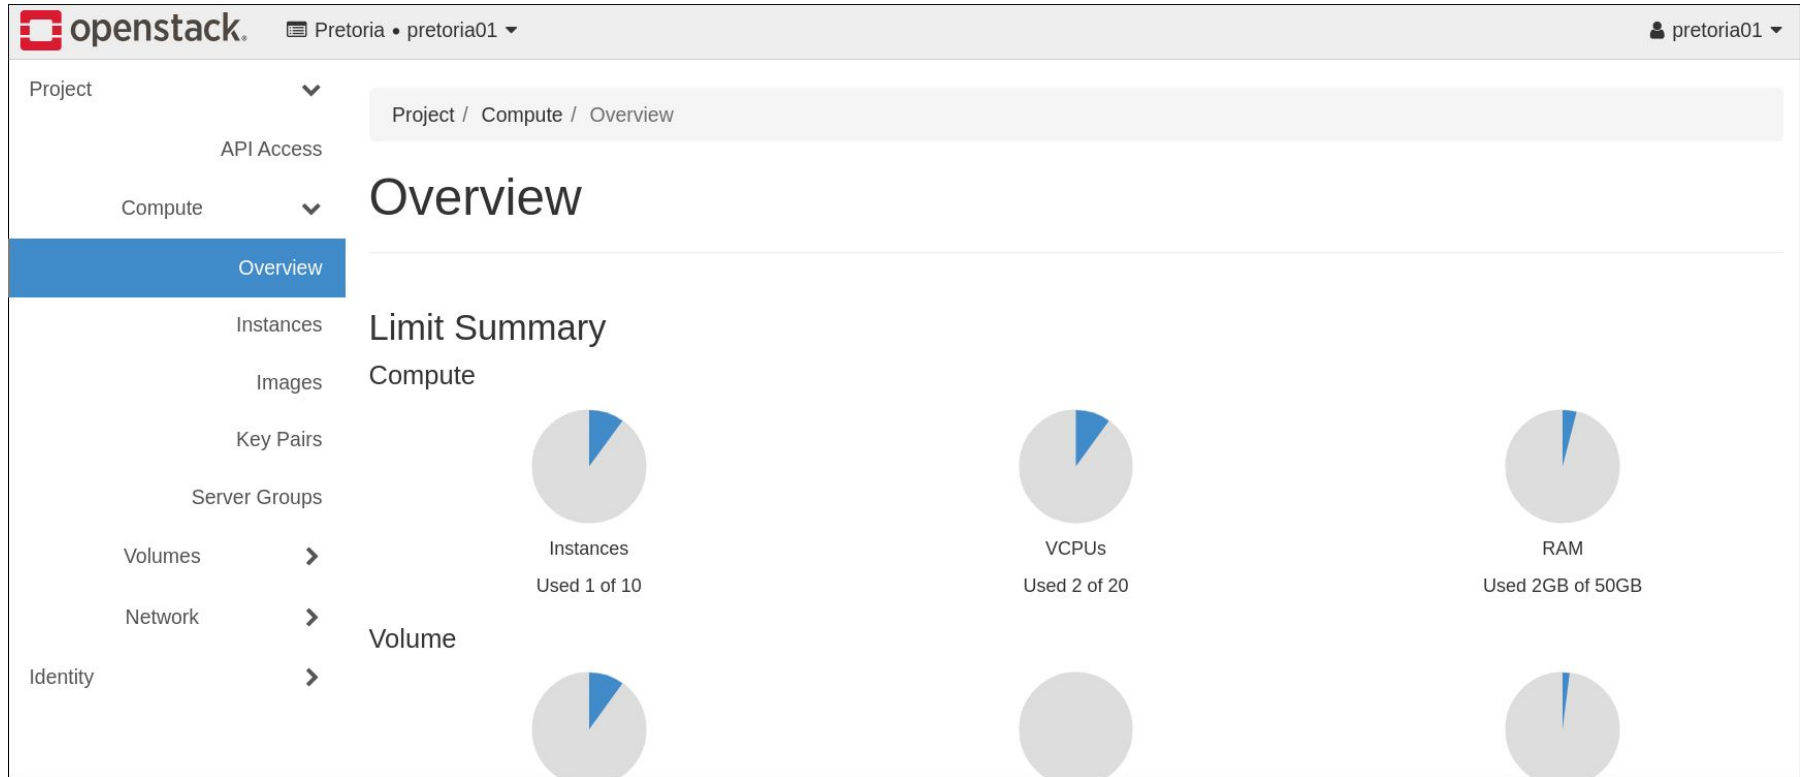
- ▷ **Image**: single file that contains a virtual disk that has a bootable operating system installed on it.
  - Used to create VM instances;
- ▷ **Volume**: block storage devices that you attach to instances to enable persistent storage.
  - You can attach or detach a volume to a running instance. You can also create a snapshot from or delete a volume.

# Openstack

## Main concepts

- ▷ **Snapshot:** provides a copy of a currently running VM or volume which can be stored into an external service such as “Glance”.
  - This can be used to “*pre-configure*” a VM, installing required software beforehand;

# Openstack - The horizon interface





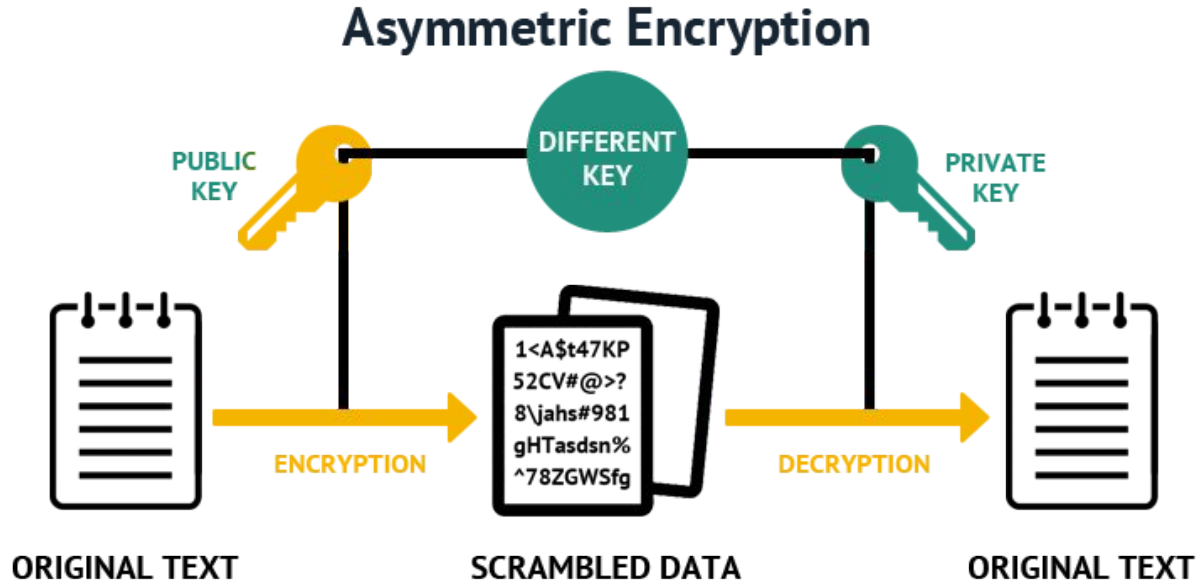
# Openstack

## VM Creation process

- ▷ Creating virtual machines needs a combination of:
  - An **image** from which the VM should be created;
  - A **flavor** defining the size of the virtual machine;
  - (optional) A **virtual network** to connect you VM to;
  - A set of firewall rules (**security groups**);
  - A **keypair** to define the authentication for logging in as the VM administrator;

## Accessing remote servers

- ▷ Private-Public key pairs;
- ▷ Asymmetric Cryptography;



# Openstack

## Installing your own Openstack

- ▷ [Packstack](#) - Cloud Proof of Concept
- ▷ All-in-one concept cloud;
- ▷ Get an Openstack cloud ready in less than 15 minutes;
  - Can run on a notebook (**considering enough resources are available**)

# Agenda

- ▷ Introduction/motivation;
- ▷ A few words about virtualization;
- ▷ Cloud environments;
  - Public vs Private vs Hybrid
  - Cloud deployment models: SaaS, PaaS, and IaaS;
- ▷ Jupyter Notebooks;
- ▷ Openstack cloud manager;
- ▷ Other virtualization techniques;

# Environment Managers

## Conda Environment Manager

- ▷ Package and **environment management system** that runs on multiple operating systems;
- ▷ Allows to **create different virtual environments**
  - **Separate software and versions** of packages/dependencies
- ▷ **Benefits:**
  - **No need for administrator access** to the computer to perform library environment modifications

# Environment Managers

## Conda Environment Manager

- ▷ Install from the [link](#);
- ▷ Create environments:

```
conda create -n ENVIRONMENT_NAME
```

- ▷ (de)Activate environment:

```
conda [de]activate ENVIRONMENT_NAME
```

- ▷ Install software to an environment:

```
conda install -n ENVIRONMENT_NAME  
SOFTWARE[=VERSION]
```

# Environment Managers

## Conda Environment Manager

- ▷ A lot more to it:
  - Separate channels for distributing softwares;
  - Ability to share your environments;
    - Export/import environments;
- ▷ Read the [Cheat Sheet](#)

# Containers

## Docker Containers

- ▷ Segregation of processes in the same operating system;
- ▷ Maximum isolation possible from the rest of the environment;
- ▷ File Systems created from an “image”;
- ▷ Makes reproducibility much easier
- ▷ Conceptually similar to virtual machines



# Containers

## Docker Containers - Advantages 😊

- ▷ Lightweight because they don't contain an operating system;
  - Have the same performance as code running on the host operating system;
- ▷ Up to three times more performance than virtual machines when running on the same hardware.
- ▷ Startup time in milliseconds (compared to minutes for a virtual machine);

# Containers

## Docker Containers - Disadvantages 🙄

- ▷ Always run on the Linux operating system (containers share the host operating system):
  - Virtual machines can run a different operating system for each virtual machine;
- ▷ Containers use process-level isolation
  - potentially less secure than fully isolated virtual machines